



STRATEGIC UX CRITIQUE · 19 MAY 2026 · WORKSPACE ARTIFACT

Goals as a setup step, or goals as the spine.

A senior-design comparison of two information architectures for Kustomer's agentic CX platform — **Prototype A**, where goals are the input to a 5-step wizard and ongoing work returns to a per-automation surface, and **Prototype B**, where goals are the permanent organizing unit and monitoring/improvement live in a rich Performance area. Includes a hybrid recommendation, not a winner-takes-all answer.

Prototype A

Goals Based AI Outcome (setup-wizard model)

Prototype B

Hybrid Goals-First Prototype (goal-as-spine)

Author

Senior UX critique, prepared for Product Owner

Scope

Onboarding · ongoing management · enterprise scale

- Summary
- IA trees
- Matrix
- Strengths
- Cognitive load
- Personas
- Hybrid
- Wireframes
- Verdict

01 Executive summary

The two prototypes resolve the same brief — make Kustomer AI easier to configure and more goal-driven — by inverting each other's bet. **A optimizes for the first 30 minutes; B optimizes for the next 12 months.** Neither is wrong; the right answer is to keep A's intake clarity and graduate users into B's outcome-spine for ongoing operations.

PROTOTYPE A

Goals as a 5-step setup wizard.

Set Goals · Build · Test & Evaluate · Deploy · Analyze sit in a stepper. After setup, work returns to a tool-first surface (Refund Order / Manage Automations / Settings).

WHERE IT WINS

- + Strongest first-run path of either design — single form, single CTA, very low cognitive load.
- + Auto-derivation cascade is real: goals × topics × audience produce procedures, monitors, scenarios.
- + Step indicator (Step 2 of 5, blue chip) gives clear progress + reassurance.

WHERE IT BREAKS DOWN

- Goals disappear from the nav once the wizard ends — users lose the frame they used to configure.
- Sub-step nav inside Set Goals (Build/Test/Deploy/Analyze) collides with the legacy top-level Build/Test/Deploy concepts.
- Manage Automations + per-automation cards remind users this is still tool-first under the surface.

- + Audience pills (AIC blue / AIR purple) consistently labelled on every artifact.
- + Setup Assistant in-place during intake reduces the "ask a human" friction.
- No home for Quality Monitors, Suggestions, Anomaly Detection, or Computed Fields from the M1 spec.
- Performance is one screen — fine for week one, thin at month six.

PROTOTYPE B (HYBRID)

Goals as the permanent spine, slide-out for config, Performance for depth.

Get Started wizard is the front door; Goals Home + a lightweight goal slide-out replace the 11-tab workspace; Performance hosts Quality Monitors, Suggestions, Anomalies, Reports — filterable by goal.

WHERE IT WINS

- + Outcomes are the organizing unit forever — matches the strategic brief past day one.
- + Goal slide-out is honest with M1: thin config, jump-links to depth — no over-promised editor.
- + Tabs in Performance map 1:1 to canonical M1 primitives (Quality Monitors, Suggestions, Anomalies, Reports).
- + Shared-resources model (Procedures, Knowledge, Tools, Guardrails) scales — goals reference, don't own.
- + Legacy preserved under Advanced with a LEGACY pill — honest migration path.
- + State-driven color (red wash for "needs attention") gives instant orientation.

WHERE IT STRAINS

- Drilling into a goal's monitors now requires a Performance jump with a goal filter — one extra hop.
- "Plan editing lives in Resources" is structurally clean but takes a moment to land for first-time admins.
- No persistent KPI hero per-goal — the slide-out is purposely sparse.
- Setup Assistant is no longer present during the wizard — only the Goals Assistant in the rail.

The short version

Prototype A is the better first hour. Prototype B is the better first year. A's wizard is a near-perfect intake — keep it. B's outcome spine is the only structure that survives a customer asking "how is my AI doing on CSAT?" six months in — keep it. Where they conflict — what happens after setup ends — graduate the user into B.

02 Information architecture, side-by-side


Same product, two trees. The difference is whether *Goals* is a leaf inside a wizard or the spine of the nav. Everything downstream falls out of that choice.

Prototype A · Goals Based AI Outcome

WIZARD MODEL

Kustomer AI


— ✨ Set Goals (wizard root)



Right rail: Setup Assistant (persistent, in-place during wizard).

Prototype B • Hybrid Goals-First

SPINE MODEL



Right rail: Goals Assistant (persistent, context-aware).

03 Comparison matrix

Eleven dimensions, scored 1–5. Where Prototype A and B differ by >1 pip the explanation is in the notes column.

DIMENSION	A	B	WINNER	WHY
Onboarding clarity first 30 minutes	●●●●●	●●●●●○	A	A's pill + visible Step 2/3/4 indicator + Setup Assistant in-place is the strongest intake of either design. B's wizard is structurally identical but loses the persistent Setup Assistant.
Long-term operational clarity month 3+ admin life	●●●○●	●●●●●●	B	A returns users to Refund Order / Manage Automations — they have to translate "improve CSAT" into "which automation". B keeps the goal frame and routes from goal → monitors.
Cognitive load density of competing actions	●●●●●○	●●●●●○	Tie	Both keep the right rail busy. A wins on intake (one form), B wins on browse (state-tinted rows scan instantly).
Discoverability of M1 primitives Quality Monitors / Suggestions / Anomalies	●●●○●	●●●●●●	B	A has no home for them — these are M1's flagship capabilities and they appear nowhere in the nav. B promotes them to Performance subnav.
Scalability can this hold 50 goals?	●●●○●	●●●●●●	B	A's per-automation cards in the nav don't scale past a handful. B's Goals Home is filterable, sortable, and stateful.
Migration / back-compat existing customers	●●●●●	●●●●●○	A	A is the tool-first model — every existing customer already uses it. B preserves it under Advanced [LEGACY] but the demotion can confuse legacy admins.
AI-product UX patterns does it feel agentic?	●●●●○	●●●●●●	B	A's Setup Assistant is excellent during onboarding but disappears later. B has a persistent context-aware Goals Assistant + plan-diff suggestions tightly integrated with the spine.
Enterprise readiness governance, audit, blast radius	●●●○●	●●●●●●	B	B's shared Resources library with "attached to N goals" chips is the right model for enterprise governance. A has no equivalent surface.
Visual hierarchy & polish typographic clarity	●●●●●○	●●●●●●	B	Both follow Customer's design system. B's state-driven color (red wash for attention) + slide-out animation feel more modern.
Trust / transparency "what is the AI doing?"	●●●●○	●●●●●●	B	B's per-goal Plan section names the 4 ingredients clearly (CAI behaviors / Rep AI behaviors / Procedures / Guardrails) with examples in quotes. A's Build view shows procedures but doesn't separate surfaces.
Configuration ease "set up the platform"	●●●●●	●●●●●●	Tie	Both share the same wizard intake. Equalized.

04 Strengths, friction, and missed opportunities

An opinionated read of what works, what hurts, and what each prototype *didn't build but could have*.

PROTOTYPE A

What works · what creates friction

WHAT WORKS WELL

- + The single-form intake — explicit "Skip & configure manually" affordance respects power users.
- + Numbered sections (1–8) make completion progress visible without a stepper.
- + The "Step 2 of 5" pill in the Build view ties wizard steps to real product surfaces.
- + Setup Assistant inline answers ("Suggest goals for my business", "What should escalate to a human?") demonstrate AI competency at the moment of need.
- + Test Categories with pass/warning/fail and last-run timestamps mirror a real CI pattern admins recognize.

MISSED OPPORTUNITIES

- Plan-diff suggestions ("we think you should tighten this guardrail") — invisible.
- Per-goal trend reporting tied to GoalProgressSnapshot — buried in a single combined view.
- No "Improve" surface — observations don't translate into actions inside the IA.

WHAT CREATES FRICTION

- Wizard sub-steps Build/Test/Deploy/Analyze are ambiguous against legacy global Build/Test/Deploy concepts — they read like two different products.
- "Refund Order" appearing as a fixed nav card hard-codes one automation example into the IA.
- Performance is a single scrollable dashboard that mixes per-goal KPIs, insights, and topic performance — overwhelming and undelineated.
- Quality Monitors / Anomalies are nowhere — admins who read the M1 spec will look and not find them.
- After the wizard, there's no obvious "next" — what does an admin do on day 8?

PROTOTYPE B

What works · what creates friction

WHAT WORKS WELL

- + Get Started is the default landing — empty-state = guided intake — and the wizard runs full-width.
- + Goals Home rows are instantly scannable: name · state pill · target · current · progress bar · suggestions count.
- + Goal slide-out resolves the "what is this goal made of?" question without committing to a workspace.
- + Plan section in the slide-out names the 4 ingredients with one-line definitions + examples in quotes.
- + "Suggest a change" CTA at the bottom of the Plan section turns the slide-out into an actionable surface.
- + Performance tabs are 1:1 with M1 canonical primitives — no invented surfaces.

WHAT CREATES FRICTION

- "Plan editing lives in Resources" requires a mental model — first-time admins may try to edit in the slide-out.
- The wizard's exit point lands on a slide-out, not a celebration screen — moment-of-success is muted.
- No persistent per-goal KPI hero — to see live numbers you must click into Performance.
- Goals nav is not expandable — to switch goals you go back to Goals Home (one extra click vs A's per-automation drill-in).

- + Resources is the shared library with "attached to N goals" — blast-radius is visible at edit time.

MISSED OPPORTUNITIES

- Setup Assistant inline copy during the wizard is gone — Goals Assistant in the rail is less integrated.
- "You're live" celebration screen could surface first-monitor reading instead of a TBD placeholder.
- Cross-goal comparison view — "show me CSAT across all goals" — not yet a destination.

05 Cognitive load analysis

Three intervals — first session, ongoing usage, optimization workflow — measured by *screens to first value, distinct concepts to hold in working memory, and destinations needed to answer a common question.*

WORKFLOW	PROTOTYPE A	PROTOTYPE B	WINNER
First-run setup "deploy my first AI"	1 form → 4 read screens → live. Excellent. Setup Assistant in-place.	Same form-based path, but landing on a slide-out vs a celebration moment. Goals Assistant in rail, not inline.	A
Ongoing health check "how's CSAT doing?"	Performance (single page, scroll through goals + topics + insights). 1 destination but high density.	Performance → Dashboard, or Goals Home → click goal → Sparkline + jump to Reports. Fewer concepts per screen.	B
Tune a procedure "refund logic is wrong"	Set Goals → Build → expand procedure → edit. Tied to one automation context.	Resources → Procedures → edit. Blast-radius chip warns of cross-goal impact.	B
Triage a failure "AI escalated something it shouldn't"	Performance dashboard → flagged conversations. Linear.	Anomalies tab in Performance (canonical M1 surface) → drill to conversation. More precise.	B
Add a goal second/third configuration	Re-enter wizard — risk of looking like first-time setup again.	"+ New goal" from Goals Home or Get Started — explicit and recurring.	B
Apply a suggestion plan-diff flow	Not present. Insights surface in Performance but no apply/dismiss UI.	Suggestions tab in Performance with confidence pill, before/after, Apply/Test/Dismiss.	B

06 Persona walkthroughs

Same task list, five user types, two prototypes. Where the success rate diverges is where the IA does or doesn't fit the mental model.

First-time admin Day 1

Just turned on Kustomer AI for an e-commerce brand. Knows support, not AI.

BEST IN A **Wizard intake** — single form, clear "Generate my AI setup" CTA, Setup Assistant answers questions inline.

BEST IN B **Empty state** — wizard auto-opens on first visit (no goals = guided intake), then routes to Goals Home with their new goal sitting on top.

VERDICT A by a hair for the very first hour; B catches up the moment they want to look at the result.

Support manager Week 2

Owens CSAT and escalation rate. Wants to know if AI is helping or hurting.

IN A Performance page — scrolls through 6 goal cards + 4 insights + topics table. Has to mentally bin the data into outcomes.

IN B Goals Home — sees 4 goals as rows, instantly orients on "Reduce reopen rate" trending wrong. Clicks → slide-out shows progress + jump to Performance with goal filter pre-applied.

VERDICT **B by a margin.** The mental model matches the role.

Ops team lead Month 1+

Configures procedures, attaches guardrails, manages the shared library.

IN A Build view inside Set Goals → procedures listed but tied to one automation context. No "this procedure is shared with 3 other automations" warning.

IN B Resources → Procedures, central library. "Attached to N goals" chip on every row. Edits broadcast cleanly.

VERDICT **B clearly.** A's tool-first model breaks down under shared-resource workflows.

Enterprise buyer Pre-sale

Evaluating Kustomer AI vs Sierra / Decagon. Cares about governance, audit, scale.

IN A Hard to demo what "monitoring my AI on real conversations" looks like — Performance shows topic data but not M1's Quality Monitors / Anomalies primitives.

IN B Performance subnav names exactly what the M1 spec promises (Quality Monitors / Suggestions / Anomalies / Reports). Resources shows shared-library governance.

VERDICT **B is the enterprise-credible IA.**

Daily power user Month 6+

Lives in Kustomer. Has 10+ goals running. Constantly tuning.

IN A Has to navigate per-automation cards in the nav. Manage Automations becomes a flat list of 20+ items.

IN B Goals Home with state filters (All / Needs attention / Running / Not yet live) handles 50 goals as cleanly as 5. Advanced surface for legacy edits.

VERDICT **B scales linearly; A doesn't.**

AI-curious exec Quarterly

Wants a one-screen "is our AI making us money?" view.

IN A Performance dashboard is closest — but mixes operational data with outcome data.

IN B Goals Home is exactly this view: every business outcome as a row with progress %.

VERDICT **B is the right exec surface.**

07 Recommended hybrid

Not "pick one." The hybrid is **Prototype B's structure + Prototype A's intake quality + restored Setup Assistant + one new destination (Improve)**. Below: what to keep from A, what to take from B as-is, and what to merge.

KEEP FROM A

- **The 5-step wizard** — Set Goals · Build · Test & Evaluate · Deploy · Analyze, with the visible step indicator pill.
- **Setup Assistant inline** during the wizard — "Suggest goals", "What should escalate?", "Draft a tone" prompts at the moment of need.
- **Audience pills** (AIC blue / AIR purple) carried through every goal, procedure, monitor, test.
- **Numbered sections (1–8)** in the wizard form — progress is visible without a stepper.
- **Test Categories with pass / warning / fail** visual treatment — the green Continue-to-Deploy gating is excellent.

TAKE FROM B AS-IS

- **Goals as the spine** — Goals Home is the landing after first run; the slide-out replaces a heavy workspace.
- **Performance subnav** with canonical M1 tabs (Quality Monitors · Suggestions · Anomalies · Reports).
- **Resources as shared library** with "attached to N goals" chips and the "Plan editing lives here" banner.
- **Settings folds AIC + AIR surfaces** — they are not top-level peers.
- **Advanced [LEGACY]** for the tool-first surface — preserved, not promoted.
- **State-driven color** — red wash for goals needing attention, restraint elsewhere.


MERGE / ADD NEW

- **Setup Assistant returns to the wizard** as inline panel (lift from A) — replaces the rail-only Goals Assistant during intake.
- **Promote "Improve" to a destination** — currently buried in Performance/Suggestions. A dedicated cross-goal Improve view ranks plan diffs by impact.
- **Per-goal "open in fullscreen"** escape hatch on the slide-out for power users who want the depth back temporarily.
- **"You're live" analyze step** shows first monitor reading or expected baseline, not "TBD".
- **Wizard exits to a fullscreen "Welcome to your goal" page** (single screen with goal hero, plan summary, first monitor) → THEN to the slide-out on close.
- **Goals nav supports a small "pinned goal" sub-list** (max 3 user-favorited) for daily power users who do want one-click goal switching.

08 Hybrid wireframes

Three text-mode sketches: (1) the wizard landing with the restored Setup Assistant inline, (2) Goals Home with pinned goals + states, (3) the per-goal slide-out with its "Improve" CTA promoted to the header.

1. Wizard with restored Setup Assistant

Rail	Nav	Wizard body (32px 40px)	Setup Assistant
K	Kustomer AI		+ Setup Assistant
 ai★	 Get Started	Tell us about your business	Hi! Ask me anything
home	 Goals	Pick a goal you want to move...	as you fill out the
inbox 1	 Performance ▼		form.
...	 Resources	① About your company	
	 Connections	[Acme Inc] [Retail & Ecom]	• Suggest goals
	 Settings	② What goals do you want to set?	for my business
	— For power users —	☒ Reduce escalations	• What should
	 Advanced [LEGACY]	☒ Improve CSAT	escalate to

		☐ Speed up rep responses ③ Where should your AI show up? ☒ AI for Customers ☒ AI for Reps ... (sections 4-8)	a human? • Draft a tone for my brand [Ask Customer AI]
		[Generate my plan]	

2. Goals Home with pinned goals + state-tinted attention

Goals · yellow gradient hero

GOALS

Your goals

[Browse templates][+ New]

Track how your AI is moving each outcome. Drill in to tune.

[All · 4] [Needs attention · 1] [Running · 3] [Not yet · 0]

24h | 7d | 30d ◀ pinned

★ PINNED

● Reduce escalations

Live

-34%

72%

3 suggs >

★ Improve CSAT

Live

+3.1

78%

2 suggs >

🚨 RED WASH for attention

● Reduce reopen rate

Attn ⚠

+4%

38%

5 suggs >

● Increase upsell 15%

Limited

+8.4%

56%

4 suggs >

3. Goal slide-out — Plan section + promoted "Improve" CTA

slide-out 520px

← Goals · Support efficiency

...

×

Reduce escalations

● Live

-34%

vs baseline

72% to -40% by Q3

sparkl

⚙ CONFIGURATION

Name · Target · Direction · Monitor

🔔 ALERTS

Notify me

● daily · email

🌟 IMPROVE (3 suggestions)

⊕ high Tighten escalation trigger

[Review →]

⊕ med Add KB article ...

📋 PLAN — what the AI is doing

💬 Customer AI behaviors · 6 active

"Answer order-tracking..."

[View all →]

🎧 Rep AI behaviors · 9 active

	Procedures · 7 attached	
	[Edit in Resources →]	
	♥ Guardrails · 5 active	
	→ Drill in	
	See monitors · See trend report	
[Archive]	[Cancel]	[Save changes]

09 Strategic verdict

One direction, three actions.

Go with Prototype B as the structural foundation.

It scales linearly with goal count, names the M1 primitives correctly, gives enterprise buyers a credible governance story, and matches the role-based mental model of every persona except the very-first-hour admin.

Action 1 · Restore the Setup Assistant inline during the wizard. This is the single biggest thing A does better. The Goals Assistant in the rail is too easy to ignore during a focused intake task. Bring back the inline panel with the same three suggestion prompts.

Action 2 · Add a "Welcome to your goal" celebration screen between Analyze (step 5) and the goal slide-out. Show the first live monitor reading or expected baseline, the plan summary, and a single "Open my goal" CTA. The wizard exit should feel like a launch, not a pop-in.

Action 3 · Promote Improve to a top-level destination, or at minimum to a permanent slot in the Performance subnav with its own badge count. Suggestions are the agent's most demonstrable intelligence — they should not be buried.

Everything else in B is correct as drawn. Ship the spine; iterate the surface.